

Problem 1

- a) First, the kernel pushes the `%eax` register onto the stack. It also pushes all of the registers that will likely be used. It allocates a stack area for each user-level thread. Then, it adds the `thread_info` struct at the lowest memory address of the stack. It masks the stack pointer to find the beginning of the `thread_info` struct. Next, the bitwise flags word is examined. After that, data validation takes place. Since the system call is negative, the `syscall_badsys` code is jumped to, which places an error code into the stack where the `eax` register was placed. Once the system call handler is exited, the value is popped into `%eax`. Since an error occurred, a negative value is returned. In this case, the value returned is `-ENOSYS`. Because of the UNIX API, the system call returns `-1`, and `errno` is set to `ENOSYS`.
- b) The handler terminates by executing the special `iret` instruction, which resets the privilege level to the previous value. This is accomplished through the `%eflags` register, which has been restored from the stack.
- c)
 - i) In a fully pre-emptive Linux kernel, once the interrupt handler completes, a context switch takes place. The execution of Process 123 is temporarily suspended and Process 999 gets the CPU. Some time later, Process 123 resumes later.
 - ii) In a non pre-emptive Linux kernel, once the interrupt handler completes, Process 123 continues execution. Then, once the system call finishes, upon return to user mode, pre-emption takes place. Now, Process 999 gets the CPU.